


Phrase 14

Table 1. FF++ Adapted Model Evaluation Results

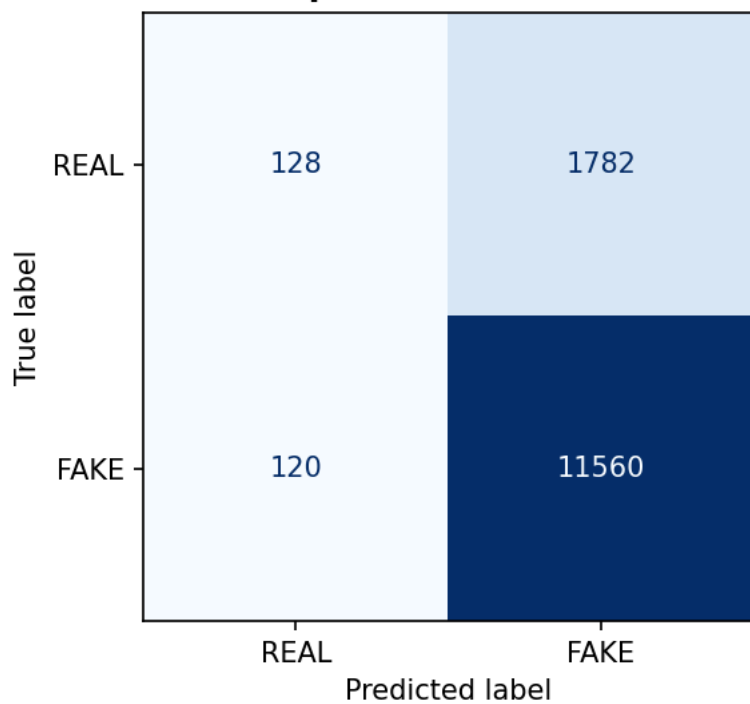
Metric	Value
Accuracy	0.8600
Precision	0.8664
Recall	0.9897
F1-Score	0.9240
Specificity	0.0670
AUC-ROC	0.7389

Table 2. Evaluation Interpretation

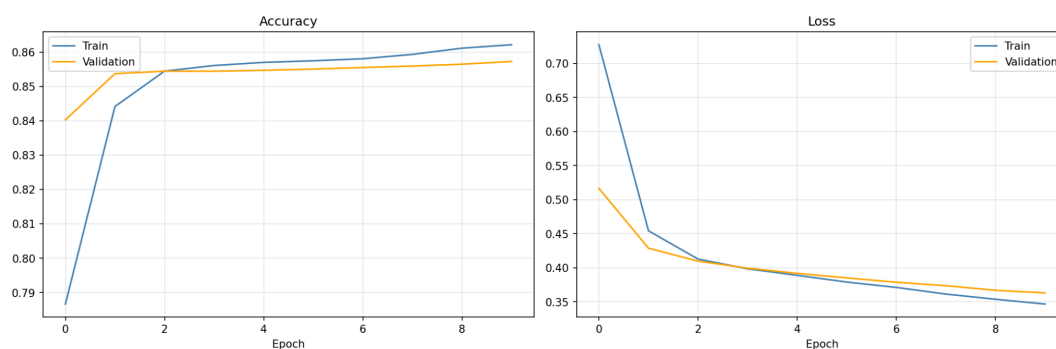
Observation	Interpretation
Accuracy	The model achieves an overall accuracy of 86.00% on the test dataset.
Precision	A precision of 86.64% indicates that most images predicted as fake are correctly classified.
Recall	The very high recall (98.97%) shows that the model detects almost all fake images.
F1-Score	The F1-score of 92.40% reflects strong performance due to the extremely high recall.
Specificity	The very low specificity (6.70%) indicates that the model misclassifies most real images as fake, resulting in a high false positive rate.
AUC-ROC	An AUC-ROC of 0.7389 indicates moderate discriminative ability between the two classes.
Key Finding	Although the model detects fake images very effectively, its poor specificity reveals a strong bias toward predicting the fake class, leading to unbalanced classification performance.



FF++ Adapted — Confusion Matrix



FF++ Domain Adaptation — Training History



Sample Extracted Frames — train

